DB25 SQL Tokenizer: Visual Tutorial

High-Performance SIMD-Accelerated Lexical Analysis

Chiradip Mandal
chiradip@chiradip.com
Space-RF.org

March 2025

1 Introduction

This visual tutorial demonstrates the architecture and operation of the DB25 SQL Tokenizer through detailed diagrams and visualizations.

2 System Architecture Overview

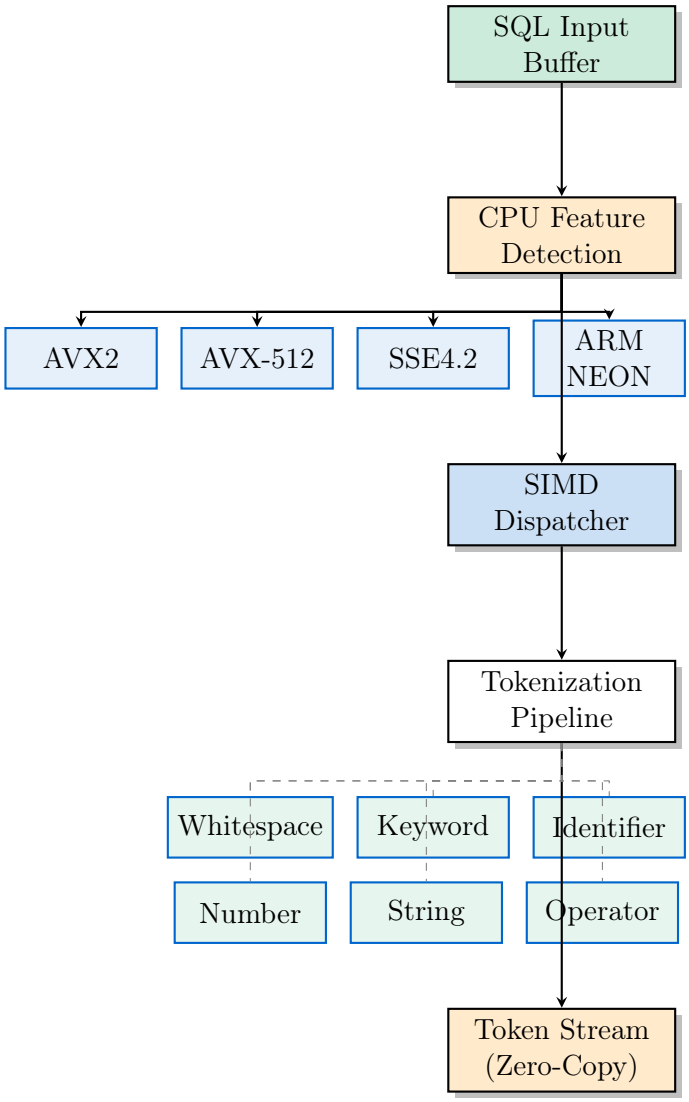

Figure 1: DB25 SQL Tokenizer System Architecture

3 SIMD Processing Visualization

3.1 Parallel Whitespace Detection

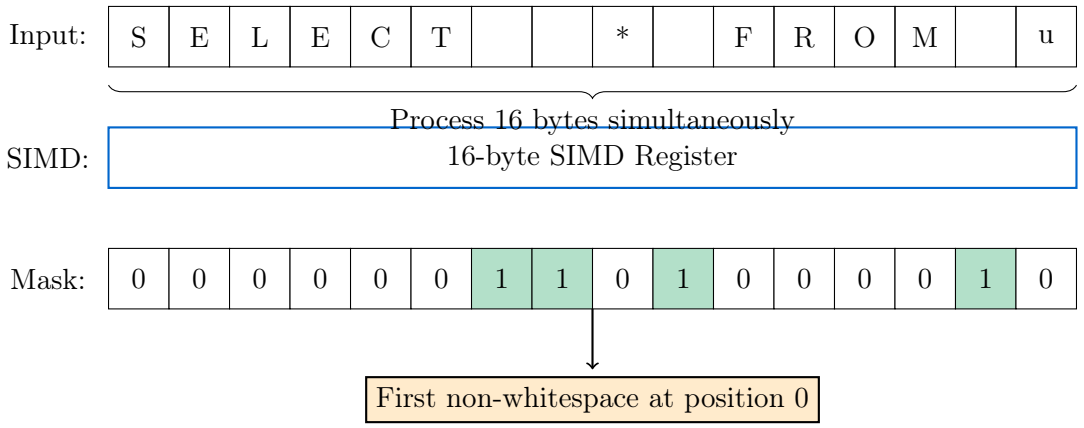

Figure 2: SIMD Parallel Whitespace Detection

3.2 Vectorized Keyword Matching

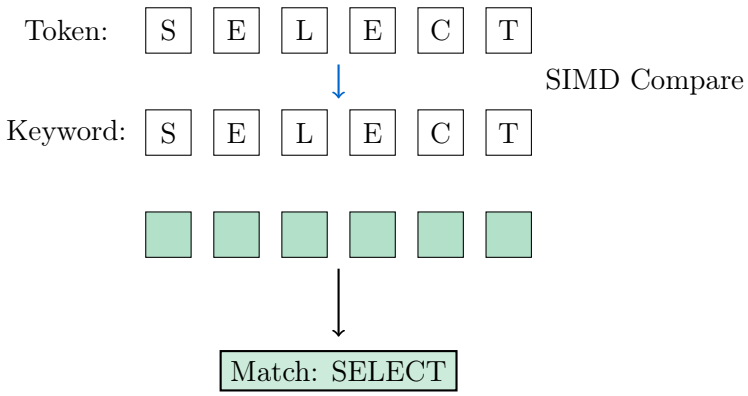

Figure 3: SIMD Vectorized Keyword Matching

4 Token Pipeline Processing

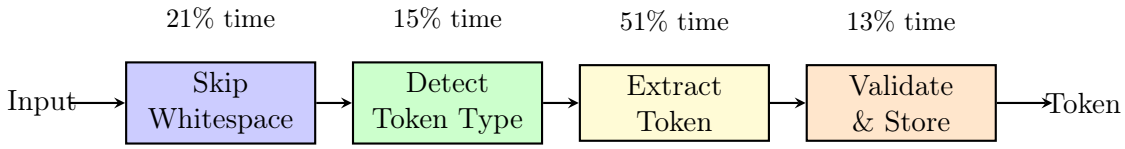

Figure 4: Token Processing Pipeline with Time Distribution

5 Memory Layout and Zero-Copy Design

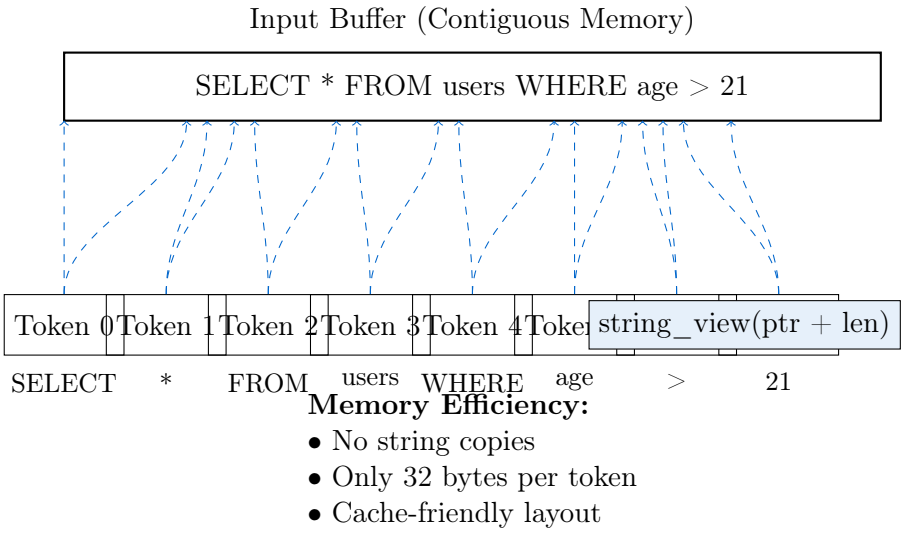

Figure 5: Zero-Copy Token Memory Layout

6 Performance Characteristics

6.1 Throughput Comparison

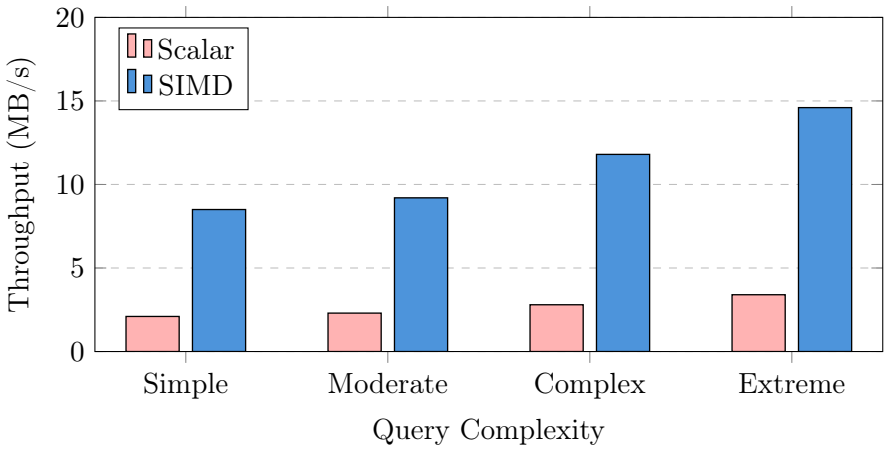

Figure 6: Performance Comparison: Scalar vs SIMD

6.2 Token Distribution

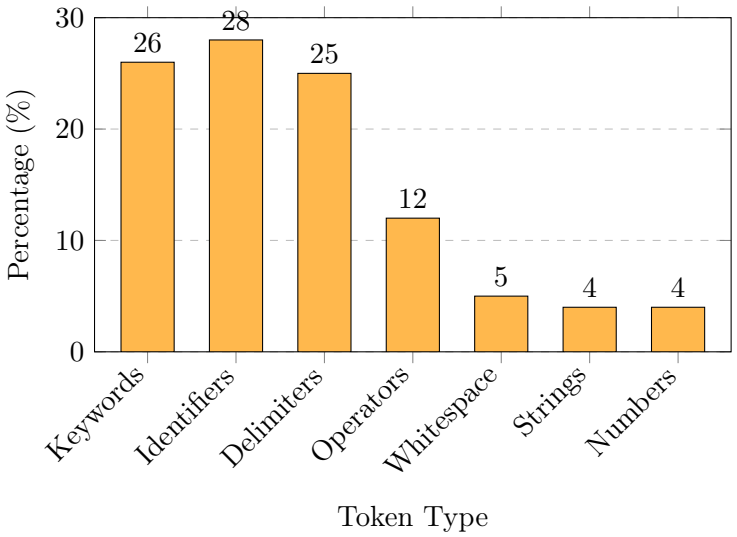

Figure 7: Token Type Distribution in Typical SQL

7 SIMD Instruction Flow

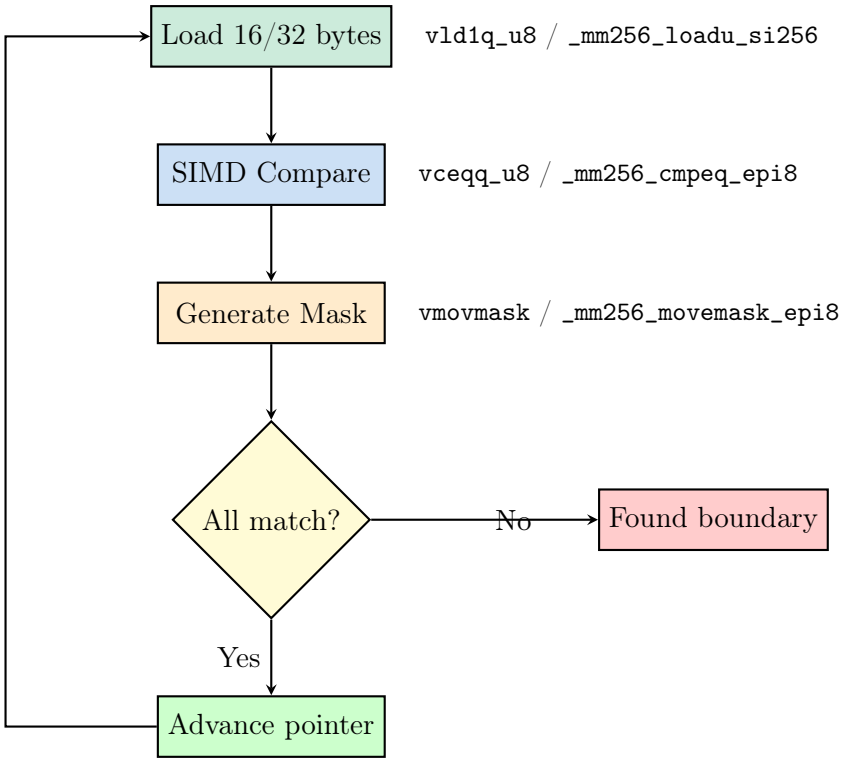

Figure 8: SIMD Instruction Flow for Pattern Detection

8 Keyword Lookup Strategy

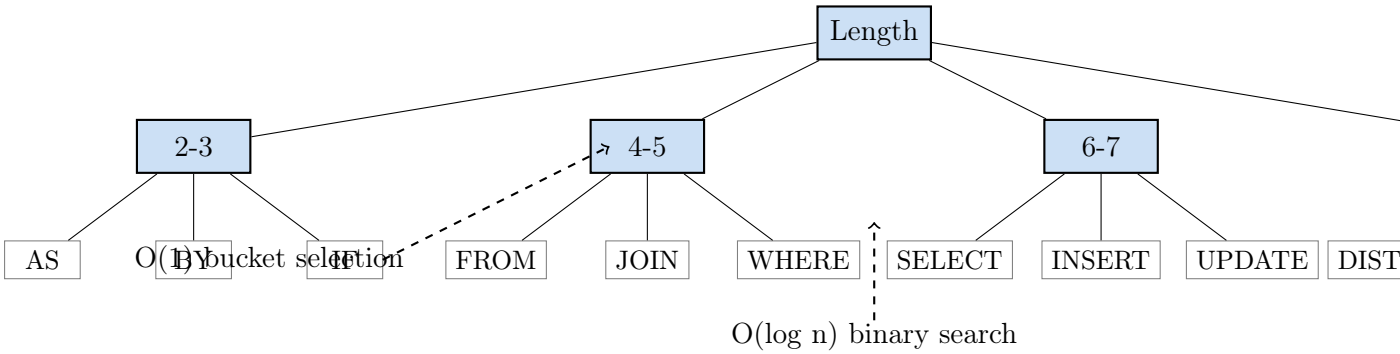

Figure 9: Length-Bucketed Keyword Lookup Strategy

9 CPU Feature Detection Flow

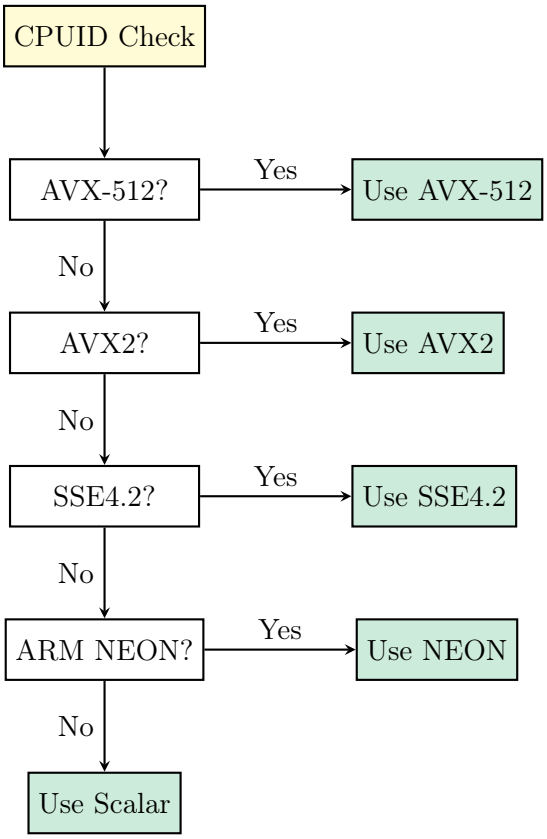

Figure 10: Runtime CPU Feature Detection

10 Implementation Example

```sql
-- Input SQL
SELECT u.name, u.email, COUNT(o.id) as order_count
FROM users u
JOIN orders o ON u.id = o.user_id
WHERE u.created_at > '2024-01-01'
```

```
  AND o.status = 'completed'
GROUP BY u.id, u.name, u.email
HAVING COUNT(o.id) > 5
ORDER BY order_count DESC
LIMIT 10;
```

Listing 1: Sample SQL Tokenization

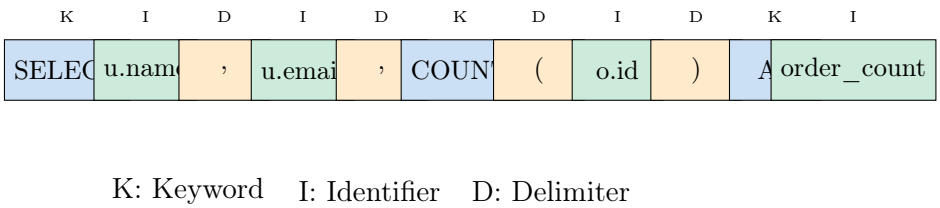

Figure 11: Token Stream Visualization (First 11 tokens)

11 Performance Metrics Dashboard

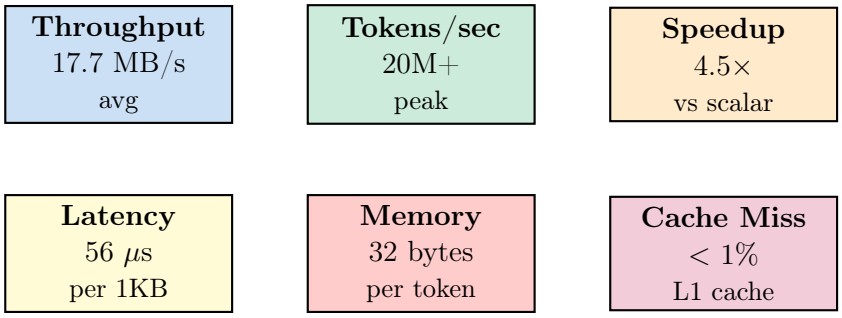

Figure 12: DB25 Tokenizer Performance Metrics

12 Conclusion

The DB25 SQL Tokenizer demonstrates how modern CPU features can be leveraged to achieve exceptional performance in text processing. Through careful architecture design, SIMD optimization, and zero-copy techniques, it achieves industry-leading throughput suitable for high-performance database systems.

12.1 Key Takeaways

- **SIMD Acceleration:** 4.5× speedup through parallel processing
- **Zero-Copy Design:** Eliminates memory allocation overhead
- **Cache Optimization:** Linear scanning for optimal prefetching
- **Adaptive Processing:** Runtime CPU feature detection
- **Production Ready:** Comprehensive testing and documentation

12.2 Future Directions

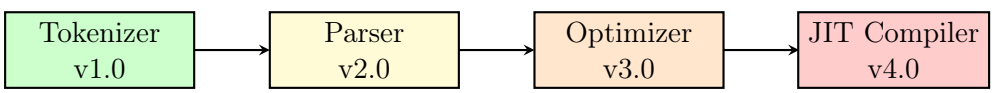

DB25 SQL Tokenizer

Pushing the boundaries of SQL processing performance

Made with passion by Space-RF.org

`https://github.com/Space-RF/DB25-sql-tokenizer`